

Introduction to Systems Programming & Linux

CST 357/457 – Systems Programming

Michael Ruth, Ph.D.

Associate Professor

Computer Science & I.T.

mruth@roosevelt.edu



Objectives

- Discuss the motivation of the course
Explain the three cornerstones of system programming including system calls, the C library, and the C compiler
- Discuss Linux programming concepts including files, directories, processes, users, groups, and file permissions

CST 357/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

System Programming?

- System Programming is the practice of writing **system software**
 - Code that lives at a **low level**
 - *Talks directly to kernel and core system libraries*
- Application v System Programming:
 - System programming requires an acute understanding of the system including its hardware and software where there are few abstractions

CST 357/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

Why Systems Programming?

- We've been trending towards application programming from systems programming
 - Web software, managed code, etc.
 - These more portable systems are still developed using systems programming
- UNIX/Linux code is almost all written at the systems level
- Finally, understanding the lower levels helps you become better in the higher levels

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

4

Cornerstones of System Programming

- System Calls
- The C Library (glibc)
- C Compiler (gcc)

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

5

System Calls

- System programming begins/ends with **system calls**
 - Function invocations made from user space into the kernel in order to request a resource
 - Linux has fewer than other systems (especially Windows which has 1000s)
 - Each architecture can implement their own as well
 - About 90% of all Linux system calls are implemented by all architectures, so that will be our focus in this class
 - User-space applications cannot directly execute kernel code or manipulate kernel data (security/reliability)
 - Kernel provides a mechanism that user-space applications “**signals**” the kernel that they wish to invoke a system call
 - The application can then **trap** into the kernel and execute only what the kernel allows it to execute
 - The user-space interacts with the kernel-space using registers

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

6

The C Library (libc)

- The C Library is at the heart of UNIX applications
 - Even when using other tools, the C language is likely underneath handling system calls for the higher-level languages
- On modern systems, the C library is provided by GNUlibc, abbreviated by glibc
 - It provides wrappers for system calls, threading support, and basic application facilities

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

7

The C Compiler (gcc)

- In Linux, the standard C compiler is the GNU Compiler Collection (gcc).
 - Originally, gcc was GNU's version of cc
 - Now, it supports many more languages than just C, hence the new name
- The compiler is very relevant to systems programming since it implements the C standard and the system ABI

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu

8

ABI v. API

- Programmers are naturally interested in portability
 - At the system level, there are two sets of definitions that impact that portability
 - API: Application Programming Interface
 - Defines the interfaces between two separate software entities communicate at the source level
 - » **Source-compatibility**
 - ABI: Application Binary Interface
 - Defines the binary interface between two components
 - » **Binary-compatibility**
 - Concerned with issues such as calling conventions, byte ordering, register use, system call invocation, linking, library behavior, and binary object format
 - Standardization is difficult: OS's tend to define their own ABIs/ABIs are tied to the architecture, so their names tend to be based on architecture
 - Enforced by the toolchain (compiler, linker, etc.), but knowledge of the ABI can lead to better code optimization

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu


9

Standards

- UNIX system programming is an old art
 - It's core hasn't changed in decades
 - However, UNIX systems are a dynamic beast
- Standards groups have formed to try to bring order to the chaos of UNIX
 - Linux doesn't comply with any of them ☺
 - Linux, instead, aims to comply with the two most important ones: Portable Operating System Interface (POSIX)/Single UNIX Specification (SUS)
 - Together, they document the C API for a UNIX-like OS interface and define system programming

C Language Standards

- For many, many years Dennis Ritchie & Brian Kernighan's Book "The C Programming Language" (written in 1978) acted as the informal specification of the language
 - Called K&R C
 - Was eventually replaced by ANSI C
 - ISO ratified it as C90
 - In 95, it was updated C95 and again in 99 (C99)
 - The most recent is C11



Concepts of Linux Programming

- All UNIX-like systems provide a mutual set of abstractions and interfaces that taken together define UNIX
 - I will cover these UNIX environmental concepts as they become important
 - However, we'll begin with some important concepts as they relate to files

Files & UNIX

- The file is the most basic and fundamental abstraction in UNIX/Linux
 - **Everything is a file!**
 - Consequently, most interactions occur via reading and writing to files even when the object is NOT what you'd consider a file
 - Dealing with a file:
 - In order to access a file, it must be opened
 - Can be opened for reading, writing, or both
 - An open file is referenced via a unique descriptor which maps from the metadata to the actual file
 - In UNIX, this is an int called the file descriptor (fd)

Regular Files

- A regular file is what you think of when you think of a file
 - Bytes of data organized into a stream
 - Any byte can be written to/read from
 - Starts at a specific byte which is referred to as an file offset (metadata maintains this)
 - Starts at zero and goes to size of int
 - The size of a file is measured in bytes and is called its length
 - **Changed via truncation operation (up or down)**
 - The kernel does NOT restrict concurrent file access

Files & File Names

- Files are usually accessed by **filenames**, they are not directly associated
 - Files are referenced by an **inode** which is assigned an integer value unique to FS
 - Value is called the **inode number**
 - Inodes store the metadata associated with files but not filenames!
 - Both an physical object & a conceptual entity

Directories

- Accessing a file by its inode number is cumbersome for humans, so files are opened from user space by name
 - Directories provide the names
 - Acts as a mapping of human-readable names to inode numbers
 - A inode and a file name pair is referred to as a **link**
 - The kernel internally uses these links to match names with inodes

File Paths & The Kernel


- Initially, there is only one directory on the disk known as the root directory
 - Denoted by /
 - How does the kernel find a directory?
 - The links inside the directory can point to the inodes of other directories
 - Which means we can have directories (concept) inside other directories (concept)!
 - The kernel walks each directory entry (called a dentry) to find the pathname of the next entry
 - Pathname resolution

Pathnames

- Every file has a pathname
 - Can be either absolute or relative

- **Absolute**


- A pathname is built by tracing a path from the root directory, through all intermediate directories, to the file
- String all the filenames in the path together, separating them with slashes (/) and preceding them with the root directory (/)
- ~(Tilde) in Pathnames:
 - the shell expands the ~ into the users home directory
- Example:
 - /home/mruth/Work1.c



Relative pathnames

- Relative
 - The pathname is built by tracing a path from the current working directory to where the file exists
 - Every directory has two intrinsic directories
 - . (this directory)
 - .. (the parent directory)
 - These make using *relative pathnames easy*

Relative Pathnames Examples



Links

- Multiple names for same file
- **Hard Link**
 - Pointer to Inode
 - Can't cross partitions
 - File removed when all links deleted
- **Symbolic (Soft) Links**
 - Pointer to file path name
 - Dangling **symlink** – *Real* file which no longer exists

Special Files

- Special (Device) File
 - A special File is a means of accessing hardware devices, including the keyboard, hard disk, CD-ROM drive, tape drive and printer
- Character Special Files
 - Correspond to character-oriented devices (e.g., Keyboard)
 - Transfer unit: byte
 - Example: /dev/console
- Block Special Files
 - Correspond to block-oriented devices (e.g., a disk)
 - Transfer Unit: group of bytes (Block)
 - Example: /dev/hda

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu
22

Questions Left Unanswered...

- File System
 - method for storing and organizing computer files and the data they contain to make it easy to find and access them
- Two basic problems:
 - Organization/Management of files
 - Mapping the file system to the storage device

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1



Michael Ruth, Ph.D.
mruth@roosevelt.edu
23

The (hierarchical) UNIX File System

- A directory is an entity in a file system which contains a group of files and/or other directories
 - A typical file system contains thousands of files, and directories help organize them by keeping related files together
- A directory contained inside another directory is called a subdirectory of that directory
- Together, the directories form a hierarchy, or tree structure
 - The root of this tree is denoted by “/”

CSF 35/457 Systems Programming
Introduction to Systems Programming
Reading: Chapter 1

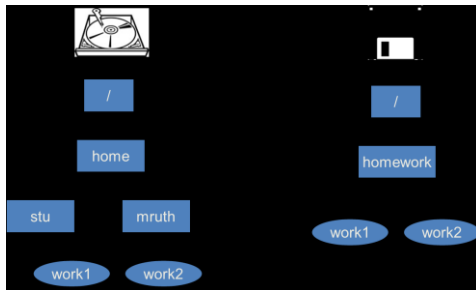


Michael Ruth, Ph.D.
mruth@roosevelt.edu
24

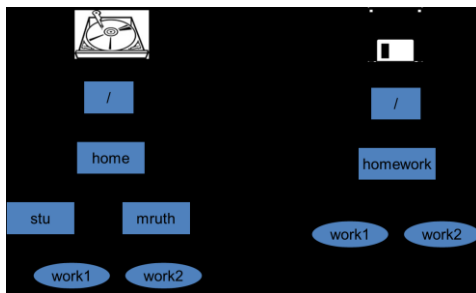
Virtual File Systems

- In UNIX, all the files on all devices appear to exist in a single hierarchy
 - Termed a Virtual FS
 - Layer of abstraction on top of a concrete FS
- ERGO, we must attach new storage onto the file system hierarchy somewhere in the hierarchy itself
 - The file system must be **mounted**
 - Mounting means “Take the file system from this CD-ROM and make it appear under this directory”
 - That directory is called the mount point

Multiple Disks and UFS



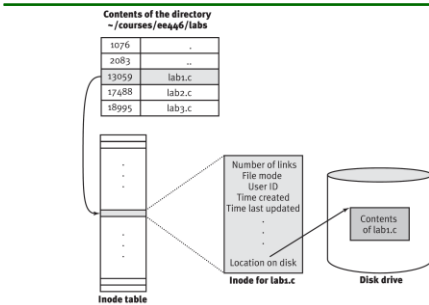
Mounting a drive (into the tree)



File Systems & Disks

- As mentioned earlier, the attributes of a file are stored in an **inode**
 - Upon creation, each file is given a unique inode from a list of disk inodes (i-list)
 - The kernel also maintains a list of inodes for open files called the inode table
 - The inode number is used to index into the table and the i-list
 - The inode contains a reference to the storage location of the file!

Pictorially...



Processes

- Processes are object code in execution: active, running programs
- A process is also associated with various resources which are managed by kernel
 - Kernel manages these resources for the processes using a process descriptor
- A process is a virtualization abstraction
 - The kernel provides each process with a virtualized processor and memory
 - The kernel handles the multiple processes, but each process believes it's on its own

Users

- Authorization in Linux is provided by users and groups
 - Each user is associated with an UID
 - Each process is associated with a real ID
 - Each process has an effective ID
 - Processes can switch users during execution
- There is a special user named **root**
 - UID is 0
 - Has special privileges that allow them to do any legal command on the system

Users & Groups

- Groups designed to group similar accounts together (simplifies administration)
- Everyone must belong to at least one group (primary group)!
 - Can belong to many supplemental groups
- Each process therefore also has a real gid, effective gid
 - Processes tend to belong to the primary

Permissions

- Each file is associated with an owning user, owning group, and three sets of permission bits
 - These bits describe the abilities of the 3
 - These things are all stored in the inode

User Type	Permission Type		
	Read (r)	Write (w)	Execute (x)
User (u)	X	X	X
Group (g)	X	X	X
Others (o)	X	X	X

Summary

- Discussed the motivation of the course
- Explained the three cornerstones of system programming including system calls, the C library, and the C compiler
- Discussed Linux programming concepts including files, directories, processes, users, groups, and file permissions

Questions?