

## Buffered I/O



CST 357/457 – Systems Programming  
Michael Ruth, Ph.D.  
Associate Professor  
Computer Science & I.T.  
mruth@roosevelt.edu

## Objectives

- Discuss buffering and user-buffering concepts including block size & kernel
- Explain the standard library components
- Discuss several buffered-I/O operations such as opening, closing, and seeking
- Explain several means to read/write files including single chars, lines, & binary
- Explain buffering control and types

CST 357/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Buffering?

- All I/O is done in **blocks**
  - So even if you write a single byte, the OS will actually write an entire block
    - This can be very inefficient as the OS has to fix up your I/O ensuring everything is block aligned
    - The entire situation is made worse by situation in which we read the same bytes over and over
- Solution: **User-buffered I/O**
  - Applications read/write naturally, but the I/O occurs in units of the file system block size

CST 357/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## User-Buffered I/O

- User-Buffered implies that the buffering happens in user space rather than the OS
  - buffering aimed to improve performance
- If we write in blocks, we can incur enormous gains from fewer writes
  - The bigger the block, the more you write, so the less often you write
  - However, our performance degrades if we are not using block boundaries
    - we're still writing partial blocks

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Block Size

- In practice, blocks are usually powers of 2
  - 512, 1024, 2048, 4096, or 8192
- We can determine a precise block size through **stat**
  - In practice, we don't really care
- Your primary goal is not to pick an oddball size (like 1130) and pick something that is a integer multiple of the actual block size
  - Since we're talking powers of 2s, this is easy...
    - Pick 4096 or 8192
    - All the smaller sizes are multiples of these
- Ok, so now we have a size... done... You wish!
  - Programs work in bytes (integers, characters, etc.), lines, etc.
  - Our user-buffer handles the difference...
    - all reads/writes go through a buffer!

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Standard Library

- The standard C library provides the standard I/O library (stdio)
  - Provides platform-independent, user-buffering
- **File pointers**
  - Stdio routines do NOT operate using file descriptors
  - Instead, the use a file pointer
    - Maps to a file descriptor
    - Represented by a pointer to the **FILE typedef**
  - Now that we're using pointes, we'll be referring to **streams**
    - Streams are a portable way of reading and writing data

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Opening a File

- **FILE \*fopen(char \*name, char \*mode)**
  - Returns a pointer to a FILE
    - Or NULL if the open failed
  - Legal values for mode include:
    - “r” – open a text file for read only
    - “w” – create text file for writing (discard previous contents)
    - “a” – open or create a text file for append operations
    - “r+” – open text file for updating (reading and writing)
    - “w+” – create text file for updating (discard old file)
    - “a+” – open or create a text file for appending operations
  - Appending a “b” to the end of any of these modes indicates the file that is being opened is a binary file

CSF 357/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Opening a file with fds

- **FILE \*fdopen(int fd, const char \*mode)**
  - Returns a pointer to a FILE
    - Or NULL if the open failed
  - The modes are the same for fopen
- NOTE:
  - Once a file descriptor is opened as a stream, it should not be used for direct I/O

CSF 357/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Closing a Stream

- **int fclose(FILE \*stream)**
  - Any buffered data is flushed
  - On success, it returns 0
    - Otherwise returns EOF and sets errno appropriately
- We can close all streams!
  - **int fcloseall(void)**
    - All streams flushed and closed
    - Always returns 0

CSF 357/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Reading from a Stream

- The standard C library provides many ways to read from a stream, but we'll focus on three simple ideas:
  - Read one character at a time
  - Read an entire line at a time
  - Reading binary data

---

---

---

---

---

---

---

---

## Reading A Single Character

- `int fgetc(FILE *stream)`
  - Returns it as a unsigned char cast as an int
    - It needs the range for error reporting

```
int c = fgetc(stream);  
if (c == EOF)  
    /* error */  
else  
    printf("c=%c\n", (char)c)
```

---

---

---

---

---

---

---

---

## Reading an Entire Line

- `char * fgets(char *s, int n, FILE *stream)`
  - Reads at most n-1 characters in the array s which is terminated by “\0” stopping if a newline is encountered
    - That newline is included in the array
  - Returns s or NULL if EOF or other error occurs

---

---

---

---

---

---

---

---

## Reading Arbitrary Strings

- Sometimes you wish to read to a delimiter other than a new line
  - We have to get back to **fgetc**

```
char *s; int c; s = str;
while (--n > 0 && (c = fgetc(stream)) != EOF)
    *s++ = c;
*s = '\0'

char *s; int c = 0; s = str;
while (--n > 0 && (c = fgetc(stream)) != EOF && (*s++ = c)
    != d);
if (c == d)
    *--s = '\0';
else {
    *s = '\0';
}
```

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
13

---

---

---

---

---

---

---

---

## Reading Binary Data

- **int fread(void \*buf, size\_t size, size\_t num, FILE \*stream)**
  - The number of elements read is returned, not the number of bytes
  - must be opened with “b” suffix

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
14

---

---

---

---

---

---

---

---

## Buffered Binary I/O Example

```
char *name = "testfile";
FILE *fp;

fp = fopen(name, "xb");

int x;

int count = 0;
int read = fread(&x, sizeof(int), 1, fp);

while (read != 0) {
    count++;
    printf("%s\t%d\n", x, read);
    read = fread(&x, sizeof(int), 1, fp);
}
printf("COUNT: %d\n", count);
fclose(fp);
```

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
15

---

---

---

---

---

---

---

---

## Writing to a Stream

- Again, there are three popular approaches to writing to streams:
  - Writing a single character
  - Writing a string of characters
  - Writing binary data

---

---

---

---

---

---

---

## Writing a Single Character

- `int fputc(int c, FILE *stream)`
  - Writes the unsigned character `c` to the stream
  - On success, it returns `c...`
    - Otherwise returns EOF

```
if (fputc('p', stream) == EOF)
    /*error */
```

---

---

---

---

---

---

---

## Writing a String of Characters

- `int fputs(const char *s, FILE *stream)`
  - Writes the entire string `s` on the stream
  - Returns non-zero if successful, or EOF for error

```
if (fputs("Michael is awesome\n", stream) == EOF)
    /*error */
```

---

---

---

---

---

---

---

## Writing Binary Data

- **int fwrite(void \*buf, size\_t size, size\_t num, FILE \*stream)**
  - Writes num elements each size bytes in length from the data pointed at buf
  - The number of elements written is returned
  - must be opened with “b” suffix

---

---

---

---

---

---

---

## Seeking?

- Seeking allows us to manipulate the current file position
  - **int fseek (FILE \*stream, long offset, int whence)**
    - Whence controls the function:
      - SEEK\_SET
        - » position = offset
      - SEEK\_CUR
        - » position = current position + offset
      - SEEK\_END
        - » Position = EOF + offset
    - If successful, returns 0, otherwise returns -1

---

---

---

---

---

---

---

## If you like convenience

- **int fsetpos(FILE \*stream, fpos\_t \*pos)**
  - Basically fseek with whence set to SEEK\_SET
- **void rewind(FILE \*stream)**
  - Basically fseek(stream, 0, SEEK\_SET)

---

---

---

---

---

---

---

## Get the File Position

- **long ftell (FILE \*stream)**
  - On success, returns current file position
  - Otherwise, returns -1 and sets errno
- We also have:
  - **int fgetpos(FILE \*stream, fpos\_t \*pos)**
    - On success, returns 0 and sets the pos to current position

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
22

---

---

---

---

---

---

---

---

## Remember to Flush

- **int fflush(FILE \*stream)**
  - If stream is NULL, all streams are flushed
  - On success, returns 0, otherwise returns EOF
- Note:
  - This only flushes user buffers to kernel
  - Not necessarily going to disk, you need to flush the kernel buffers using fsync() to do that

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
23

---

---

---

---

---

---

---

---

## Errors & EOF

- Some (most) of the standard I/O interfaces discussed do a poor job of communicating failures
  - Is it EOF or a literal failure?
  - Standard I/O provides two interfaces to determine whether error or EOF
    - **int feof(FILE \*stream)**
      - If EOF has been encountered, it returns nonzero and a 0 otherwise
    - **int ferror(FILE \*stream)**
      - If error indicator has been set, return nonzero, 0 otherwise

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
24

---

---

---

---

---

---

---

---



## Obtaining the File Descriptor

- Mixing standard I/O calls with system calls is not normally advised!
  - Buffering is happening on user and kernel side... so you should flush before doing so
  - Generally speaking though, this should never happen
- `int fileno(FILE *stream)`
  - Returns the fd on success, -1 otherwise

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
26

---

---

---

---

---

---

---

## Controlling the Buffering

- User-buffered I/O is based on buffering in user-space, so we can control it
  - There are three types:
    - Unbuffered
    - Line-Buffered
    - Block-Buffered
  - Usually, the default buffering type is generally optimal
    - Can be set with `setvbuf` function
      - `_IONBF`
      - `_IOLBF`
      - `_IOFBF`

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
26

---

---

---

---

---

---

---

## Summary

- Discussed buffering and user-buffering concepts including block size & kernel
- Explained the standard library components
- Discussed several buffered-I/O operations such as opening, closing, and seeking
- Explained several means to read/write files including single chars, lines, & binary
- Explained buffering control and types

CSF 35/457 Systems Programming  
Buffered I/O  
Reading: Chapter 3



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
27

---

---

---

---

---

---

---

## Questions?



CSF 357/457 Systems Programming  
Buffered I/O  
Roadmap - Chapter 3

**R** ROOSEVELT  
UNIVERSITY

Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
76

---

---

---

---

---

---

---

---