

## File System I/O



CST 357/457 – Systems Programming  
Michael Ruth, Ph.D.  
Associate Professor  
Computer Science & I.T.  
mruth@roosevelt.edu

## Objectives

- Discuss file system IO including file metadata, stat, permissions & ownership
- Discuss directory including entries, creation, removal, opening directory streams, and reading directories
- Explain links including hard links and symbolic links and their system calls
- Discuss copying and moving files

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## Files & Their Metadata

- Recall that each file is referenced by an **inode** addressed by an **inode number**
  - The inode is both a physical object located on the disk and a logical entity complete with a data structure in the kernel
  - The inode stores the **metadata** associated with a file
    - You can obtain the inode number using `-i` flag with `ls` command

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

## The stat family

- UNIX provides a family of functions to obtain the metadata of a file:
  - #include <sys/types.h>
  - #include <sys/stat.h>
  - #include <unistd.h>
  - int stat(const char \*path, struct stat \*buff)
  - int fstat(int fd, struct stat \*buff)
  - int lstat(const char \*path, struct stat \*buff)

## struct stat

```
struct stat {  
    dev_t st_dev;           /* Device ID containing file */  
    ino_t st_ino;           /* inode number */  
    mode_t st_mode;        /* permissions */  
    nlink_t st_nlink;      /* number of hard links */  
    uid_t st_uid;          /* user ID of owner */  
    gid_t st_gid;          /* group ID of owner */  
    dev_t st_rdev;         /* device ID (special file) */  
    off_t st_size;         /* total size in bytes */  
    blksize_t st_blksize;  /* block size for file system */  
    blkcnt_t st_blocks;    /* number of blocks */  
    time_t st_atime;       /* last access time */  
    time_t st_mtime;       /* last modification time */  
    time_t st_ctime;       /* last status change time */  
};
```

## stat Ex

```
#include <sys/types.h>  
#include <sys/stat.h>  
#include <unistd.h>  
  
int main(int argc, char *argv[]) {  
    struct stat sb;  
    int ret;  
  
    if (argc < 2) { /* error */ }  
  
    ret = stat(argv[1], &sb);  
  
    if (ret) { /* error */ }  
  
    printf("%s is %ld bytes\n", argv[1], sb.st_size);  
}
```

## stat Ex 2

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    struct stat sb;
    int ret;

    if (argc < 2) { /* error */ }

    ret = stat(argv[1], &sb);

    if (ret) { /* error */ }

    printf("File Type: ");
```

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

7

## Stat Ex 2 (cont)

```
switch(sb.st_mode & S_IFMT) {
    case S_IFBLK:
        printf("block device node \n");
        break;
    case S_IFCHR:
        printf("character device node \n");
        break;
    case S_IFDIR:
        printf("Directory\n");
        break;
    case S_IFIFO:
        printf("FIFO\n");
        break;
    case S_IFLNK:
        printf("Symbolic Link\n");
        break;
    case S_IFREG:
        printf("Regular File\n");
        break;
    case S_IFSOCK:
        printf("Socket\n");
        break;
    default:
        printf("Unknown\n");
        break;
}
```

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

8

## Permissions

- Stat can be used to obtain the permissions, two other calls set those values

```
- #include <sys/types.h>
- #include <sys/stat.h>

- int chmod(const char *path, mode_t mode)
- int fchmod(int fd, mode_t mode)
```

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu

9

## Ownership

- In stat, the st\_uid and st\_gid fields provide the file's owner/group, we can set them using the following:

```
- #include <sys/types.h>
- #include <unistd.h>

- int chown(const char *path, uid_t owner, gid_t group)
- int lchown(const char *path, uid_t owner, gid_t group)
- int fchown(int fd, uid_t owner, gid_t group)
```

---

---

---

---

---

---

---

---

## Directories

- In UNIX, a directory is a simple file consisting of a list of directory entries
  - Each entry is a inode-name mapping AKA Link
  - Every directory contains two special directories:
    - . -> right here
    - .. -> parent (except for root directory)

---

---

---

---

---

---

---

---

## Current Working Directory

- Every process has a current directory which is called its current working directory (cwd)
  - Starting point for resolving relative pathnames
  - A process can obtain & change its cwd

---

---

---

---

---

---

---

---

## Obtaining/Changing CWD

```
#include <unistd.h>
```

```
char * getcwd(char *buff, size_t  
size);
```

```
int chdir(const char *path);  
int fchdir(int fd);
```

- We often get CWD so we can return to it later...
  - If this is the case, use the fd as it is more efficient

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
13

---

---

---

---

---

---

---

---

## Creating Directories

```
#include <sys/stat.h>
```

```
#include <sys/types.h>
```

```
int mkdir(const char *path,  
mode_t mode)
```

**Final perm:**

```
(mode & ~umask & 01777)
```

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
14

---

---

---

---

---

---

---

---

## Removing Directories

```
#include <unistd.h>
```

```
int rmdir(const char *path);
```

**NOTE:**

**there is no recursive  
version!**

CST 357/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
15

---

---

---

---

---

---

---

---

## Reading a Directory's Contents

- To read a directory's contents, you need a **directory stream** represented by **DIR**

```
–#include <sys/types.h>
–#include <dirent.h>
–DIR * opendir(const char
    *name)
```

---

---

---

---

---

---

---

## Directory Stream

- A directory stream is basically a file descriptor with some metadata
  - We can obtain the fd if necessary

```
–#define _BSD_SOURCE /*or
    _SVID_SOURCE */
–#include <sys/types.h>
–#include <dirent.h>

–int dirfd(DIR *dir);
```

---

---

---

---

---

---

---

## Reading a Directory's Contents (2)

- Once you have a directory stream, you can read the directory entries:

```
#include <sys/types.h>
#include <dirent.h>

struct dirent * readdir(DIR *dir)

struct dirent {
    ino_t d_ino; /* inode number */
    off_t d_off; /* offset to next dirent */
    unsigned short d_reclen; /* length of record */
    unsigned char d_type; /* type of file */
    char d_name[256]; /* filename */
}
```

---

---

---

---

---

---

---

## Closing a Directory Stream

```
#include <sys/types.h>
#include <dirent.h>

int closedir(DIR *dir)
```

---

---

---

---

---

---

---

## Links

- A link is essentially a name for an inode
  - We can create more than one link to the same inode
    - Only on the same file system though
    - We call these **hard links**
  - Another type of link is not a filesystem mapping, but a higher level pointer than is interpreted at runtime
    - We call these **symbolic links**
    - May be relative or absolute
    - May point to nonexistent files (**dangling symlink**)

---

---

---

---

---

---

---

## Hard Links

```
#include <unistd.h>

int link(const char *oldpath,
const char *newpath)
```

---

---

---

---

---

---

---

## Symbolic Links

```
#include <unistd.h>
```

```
int symlink(const char  
*oldpath, const char  
*newpath)
```

CS1 35/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
22

---

---

---

---

---

---

---

## Unlinking & Remove

```
#include <unistd.h>
```

```
int unlink(const char *path)
```

```
#include <stdio.h>
```

```
int remove(const char *path);
```

CS1 35/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
23

---

---

---

---

---

---

---

## Copying & Moving Files

- There is no system call for copying a file
  - We just open the old file, open the new file, copy the file byte by byte, and close both
- For moving, there is a system call:
  - #include <stdio.h>
  - int rename(const char \*old, const char \*new);

CS1 35/457 Systems Programming  
File System I/O  
Reading: Chapter 8



Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
24

---

---

---

---

---

---

---



## Summary

- Discussed file system IO including file metadata, stat, permissions & ownership
- Discussed directory including entries, creation, removal, opening directory streams, and reading directories
- Explained links including hard links and symbolic links and their system calls
- Discussed copying and moving files

CST 357/457 Systems Programming  
File System I/O  
Readings: Chapter 8

**R** ROOSEVELT  
UNIVERSITY

Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
26

---

---

---

---

---

---

---

## Questions?



CST 357/457 Systems Programming  
File System I/O  
Readings: Chapter 8

**R** ROOSEVELT  
UNIVERSITY

Michael Ruth, Ph.D.  
mruth@roosevelt.edu  
26

---

---

---

---

---

---

---