

Process Management

CST 357/457 – Systems Programming
Michael Ruth, Ph.D.
Associate Professor
Computer Science & I.T.
mruth@roosevelt.edu



Objectives

- Discuss process fundamentals including process concepts, process components, APIs, and executing new processes
- Explain the process termination concepts
- Discuss managing users and groups of processes including queries and change
- Explain the daemon concept and its core implementation in a Linux & C system

CST 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu

Processes

- Processes are object code in execution: active, running programs
- A process is also associated with various resources which are managed by kernel
 - Kernel manages these resources for the processes using a process descriptor
- A process is a virtualization abstraction
 - The kernel provides each process with a virtualized processor and memory
 - The kernel handles the multiple processes, but each process believes it's on its own

CST 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu

Programs, Processes, & Threads

- A **binary** is compiled, executable code that is dormant...
 - We colloquially use the term **program**
- A **process** is a running program
 - Includes the binary, instance of virtualized memory, kernel resources, security context, and one or more threads
- A **thread** is the unit of activity inside a process
 - Has its own virtualized processor which includes a stack, processor state such as registers and an IP

CSF 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu

6

The Process ID (pid)

- Each process is represented by a unique identifier (**pid**)
 - Guaranteed to be unique at any single point in time
 - Kernel allocates PIDs in strictly linear fashion until it reaches its max (max = 32768)
 - Will not reuse PIDs until it wraps around
 - The idle process which is the process that the kernel runs when there are no other runnable process has pid 0
 - first process started is called **init** and has pid 1
 - Boots the system & launches login programs

CSF 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu

5

Process Hierarchy

- The process that spawns a new process is known as the **parent**
 - New process is called the **child**
- Every process is spawned from another process
 - So every process has a parent (**ppid**)
- Each process is **owned by a user & group**
 - Controls access rights for the process
 - Each child process inherits their user/group
- Each process is also part of a **process group**
 - Also inherited
 - Makes it easy to send signals to a group of processes

CSF 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu

6

Get the PID/PPID

```
#include <sys/types.h>
#include <unistd.h>

pid_t getpid(void);
pid_t getppid(void);
```

Running a New Process

- There are two ways to do this:
 - **exec**
 - Loads the binary program which replaces the parent
 - Called *executing* a program
 - There is no single function, there are many
 - **fork**
 - Creates a near-duplicate of the parent (initially)
 - Called *forking*

execl

```
#include <unistd.h>

int execl (const char *path,
const char *arg, ...);
```

- The ... is variadic!
 - more arguments can follow, but should be NULL terminated

exec1 EX

```
int ret;  
  
ret =  
exec1("/bin/vi", "vi", "/home/student/homework.txt", NULL)
```

What's Different/Same?

- Difference?
 - Any pending signals are lost
 - Any signals the process is catching are returned to their default behavior
 - Any memory locks are dropped
 - Most thread attributes are returned to defaults
 - Anything related to process's memory address space is cleared
 - Anything that exists in user space is cleared
- Same?
 - PID, PPID, priority, and owning user/group

The rest of the **exec** family

```
#include <unistd.h>  
  
int execlp(const char *file, const char *arg, ...);  
  
int execl(const char *path, const char *arg, ...,  
char * const envp[]);  
  
int execlp(const char *path, const char *argv[]);  
  
int execlp(const char *file, const char *argv[]);  
  
int execl(const char *file, const char *argv[],  
const char *envp[]);
```

The `fork()` System Call

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void)
```

What's Different/Same? (2)

- Difference?
 - The pid of the child is newly allocated and different from that of the parent
 - The child's parent pid is set to the pid of its parent process
 - Resource statistics are reset to zero in the child
 - Any pending signals are cleared and are not inherited by the child
 - Any acquired file locks are not inherited by the child

Common Usage of Fork

```
pid = fork();
```

```
If (!pid) {  
    ret =  
    execl("/bin/vi", "vi", "/home/student/homework.txt",  
    NULL)
```

Copy-On-Write (COW)

- Early systems duplicated everything upon the fork which is time-consuming
 - Lazy optimization strategy (COW) with a simple premise:
 - If we're all reading, no copy needs to be made
 - If we're writing, make the copy

Terminating a Process

```
#include <stdlib.h>

void exit(int status)
```

Other Ways to Terminate

- The classic way to terminate a program is simply to let it finish
 - “falling off the end”
 - implicit `exit()` call
- Good Practice:
 - Either use `exit()` or have `main` return
- A process can also terminate:
 - Process is sent the signal `SIGTERM/SIGKILL`
 - Incurring the wrath of the kernel (faults)

atexit()

- You can register a function to execute upon termination
 - Think cleanup type of operation
 - Functions execute in the reverse order they are registered (LIFO)
 - Functions cannot call exit()
- `#include <stdlib.h>`
- `int atexit(void (*function)(void));`
 - Prototype:
 - `void function(void);`

CSF 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu
16

Waiting for Terminated Children

- When a process terminates, the kernel sends SIGCHLD to the parent
 - We'll learn about signals later
- When a child terminates before its parent, it becomes a **zombie** process
 - Waits for parent to inquire about its status
 - Once parent inquires, the child ceases to exist

CSF 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu
20

Wait() & Macros

```
#include <sys/types.h>
#include <sys/wait.h>
```

```
pid_t wait(int *status)
```

- Returns the PID of the terminated child or -1 on error
- If not NULL, the status pointer contains additional information (uses MACROS)
 - `int WIFEXITED(status)` -NORMAL
 - `int WIFSIGNALED(status)` -KILLED BY SIGNAL
 - `int WIFSTOPPED(status)` -STOPPED
 - `int WIFCONTINUED(status)` -CONTINUED
 - `int WEXITSTATUS(status)`
 - `int WTERMSIG(status)`
 - `int WSTOPSIG(status)`
 - `int WCOREDUMP(status)`

CSF 357/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu
21

Launching & Waiting

```
#define _XOPEN_SOURCE
#include <stdlib.h>

int system(const char *cmd);
```

Users & Groups

- Best practices in software development encourage the use of *least-privilege*
 - Processes should execute with the minimum level of rights possible
 - IE, don't run everything as root!
- Sometime, though, we need higher privileges, so we need to change users/groups, etc

Real, Effective, & Saved UID/GIDs

- There are 4 user IDs and 4 group IDs associated with each process:
 - Real ID:
 - UID of the user who ran the process
 - Set to parent's Real UID and doesn't change
 - Effective ID:
 - UID that the process is currently wielding
 - **Permission checks happen against this value**
 - Initially the same as Real ID
 - By executing a setuid binary (suid), the process can change its effective ID
 - More exactly, it is changed to the owner of the executable file (**fileys ID**)
 - Saved ID:
 - Process's original effective UID
 - Upon an exec call, the effective UID is set to Saved ID

Changing Real/Saved UID/GID

```
#include <sys/types.h>
#include <unistd.h>
```

```
int setuid(uid_t uid);
int setgid(gid_t gid);
```

- Sets the effective UID/GID
 - if current effective UID/GID is root, then the saved and real are also set
- NOTE:
 - Root can use any value (valid UID/GID)
 - NonRoot must use either real or saved UID/GID

Changing Effective UID/GID

```
#include <sys/types.h>
#include <unistd.h>
```

```
int seteuid(uid_t uid);
int setegid(gid_t gid);
```

- Sets the effective UID/GID
- NOTE:
 - Root can use any value (valid UID/GID)
 - NonRoot must use either real or saved UID/GID

Obtaining UID/GIDs

```
#include <sys/types.h>
#include <unistd.h>
/* real */
uid_t getuid(void);
gid_t getgid(void);
/* effective */
uid_t geteuid(void);
gid_t getegid(void);
```

Sessions & Process Groups

- Each process is a member of a **process group**, which is one or more processes associated with another for job control
 - Primary attribute: send signals to entire group at once
 - each group has a **process group leader**
- A **session** is a collection of one or more process groups (usually tied to a shell)
 - Exist to consolidate logins around terminals
 - Process groups in a session are divided:
 - One foreground process group
 - Zero or more background groups
 - Each session has a **session leader**

CSF 35/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu
76

Daemons

- A daemon is a process that runs in the background
 - Normally started at boot and are run as a special user (root, apache, etc)
 - Handle system-level tasks
 - As a convention, name usually ends in a d
 - In general, two requirements:
 - Runs as child of init
 - Must NOT be connected to a terminal

CSF 35/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu
77

Steps to be a Daemon

- Call Fork()
- The parent calls exit()
 - Cleans up parent & ensures parent is init
- Call setsid()
 - Daemons want nothing to do with process groups or sessions, so it's necessary to create your own
- Changing working directory via chdir()
 - Don't wish to lock anything by accident
- Close all file descriptors
 - Don't want any inheritance
- Open file descriptors 0, 1, and 2 and redirect them to /dev/null

CSF 35/457 Systems Programming
Process Management
Reading: Chapter 5



Michael Ruth, Ph.D.
mruth@roosevelt.edu
78

Most UNIX Systems

```
#include <unistd.h>
```

```
int daemon(int nochdir, int  
noclose)
```

Summary

- Discussed process fundamentals including process concepts, process components, APIs, and executing new processes
- Explained the process termination concepts
- Discussed managing users and groups of processes including queries and change
- Explained the daemon concept and its core implementation in a Linux & C system

Questions?


