

Introduction to Threading

CST 357/457 – Systems Programming

Michael Ruth, Ph.D.

Associate Professor

Computer Science & I.T.

mruth@roosevelt.edu



Objectives

- Discuss thread concepts including terminology, benefits, and costs
- Explain threading models and patterns
- Discuss concurrency and parallelism
- Explain race conditions and synchronization using mutual exclusion
- Discuss deadlock and its avoidance
- Explain pthreads API including thread implementations and using the API

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

Programs, Processes, & Threads

- A **binary** is compiled, executable code that is dormant...
 - We colloquially use the term **program**
- A **process** is a running program
 - Includes the binary, instance of virtualized memory, kernel resources, security context, and one or more threads
- A **thread** is the unit of activity inside a process
 - Has its own virtualized processor which includes a stack, processor state such as registers and an IP
 - Smallest unit of execution schedulable by an operating system's process scheduler

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

Introduction

- Each process contains one or more threads
 - If you only have one thread, there is only a single unit of execution in the process and only one thing going on at once
 - Single-threaded processes (what you're used to)
 - If you have more than one thread, then there is more than one thing going at once
 - Multi-threaded process

CSF 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

4

OS Abstractions

- Modern OS provide two fundamental virtualized abstractions to userspace:
 - Virtualized memory
 - Associated with each process
 - Virtualized processor
 - Associated with each thread
- Together they give the illusion that each process is the only running process

CSF 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

5

Multithreading Benefits

- Programming Abstraction
 - Natural approach to some problems
- Parallelism
 - Improves throughput
- Improving Responsiveness
 - Long running processes block UI events
- Blocking I/O
- Context Switching
 - T-T switching is cheaper than P-P switching
- Memory Savings
 - Shared memory

CSF 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

6

Multithreading Costs

- Programming with threads is considerably more difficult than without
 - The general problem?
 - Concurrency with shared memory
 - How do we solve?
 - Not just **synchronization** alone!
 - Your threading model and synchronization strategies **MUST** be a part of your design

CSF 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

7

Threading Models

- Approaches to implement threads depend on where the functionality is provided: user or kernel space
 - **1:1 threading or Kernel-level threading**
 - There is a 1 to 1 relationship between what the kernel provides and what is consumed
 - **N:1 threading or user-level threading**
 - User space is where threading is implemented
 - Requires little or NO support from kernel but requires significant user-space code...
 - Benefit is that context switching is cheaper...
 - **N:M threading or hybrid threading**
 - Kernel and user threads are used and mapped N to M

CSF 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

8

Threading Patterns

- The first step in developing a threading strategy is to consider threading patterns
 - There are a myriad of abstractions and implementation details to consider but there really are only two core patterns:
 - **Thread-Per-Connection**
 - Unit of work is assigned to a thread and that thread is responsible only for that unit of work during its execution
 - **Event-Driven**
 - Since so much of threading is simply waiting on IO, we're going to decouple that waiting from the threads
 - Instead, we'll issue all IO asynchronously and use multiplexed IO to manage the flow of control
 - After callback, the multiplexed IO event loops hands the callback to a waiting thread

CSF 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu

9

Concurrency & Parallelism

- **Concurrency** is the ability of two or more threads to execute in overlapping time periods
 - Can occur without parallelism
 - Programming pattern (a way to approach problems)
- **Parallelism** is the ability to execute two or more threads simultaneously
 - Parallelism is a specific form of concurrency requiring hardware that allows for two threads to operate at the same time
 - Hardware feature achievable through concurrency

Race Conditions

- It is concurrency that introduces most of the **pain** having to do with threading
- Since threads can overlap their execution, it becomes **nondeterministic**
 - Threads share resources so accessing memory is like a “race” where what happens depends on who **gets there first**
 - Formally, a **race condition** is a situation in which the unsynchronized access of a shared resource by two or more threads leads to **erroneous program behavior**
 - the shared resource can be anything: data, hardware, kernel resources, etc
 - Data is the most common form and called a data race

Race Conditions (cont)

- The window in which a race can occur is referred to as a **critical region**
 - This is the area that needs **synchronization**
- Races are eliminated by **synchronizing threads** access to critical regions

Critical Region Example

```
double withdraw(double amt) {  
    if (balance > amt) {  
        balance = balance - amt;  
        return amt;  
    } else {  
        return -1;  
    }  
}
```

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
13

More Critical Region Examples

x++;

++x;

x = x + 4;

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
14

Synchronization

- The fundamental source of races is the critical regions are a window during which correct program behavior requires that threads **do not interleave execution**
 - We must synchronize to ensure that each thread has **mutually exclusive access** to the critical region
 - We refer to operations as **atomic** if it is indivisible
 - To the rest of the program atomic operations appear to occur instantaneously

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
15

Mutexes

- The most common mechanism for synchronization is the lock
 - Since we are enforcing mutual exclusion, we often refer to these locks as *mutexes*
- **NOTE:**
 - Nothing is magical about locks, nothing physically enforces the mutual exclusion
 - Locks are basically a *gentleman's agreement*

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7

 ROOSEVELT
UNIVERSITY

Michael Ruth, Ph.D.
mruth@roosevelt.edu
16

Critical Region Example P2

```
double withdraw(double amt) {  
    lock();  
    if (balance > amt) {  
        balance = balance - amt;  
        unlock();  
        return amt;  
    } else {  
        unlock();  
        return -1;  
    }  
}
```

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7

 ROOSEVELT
UNIVERSITY

Michael Ruth, Ph.D.
mruth@roosevelt.edu
17

Lock Data, Not Code

- As we move further in, don't forget that race conditions are based on **data** and it's manipulation by many threads
- Although we lock critical regions of code, we need to ensure that we focus on locking the **data**
 - Doing so ensures that the data won't be manipulated somewhere without proper consideration for concurrency

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7

 ROOSEVELT
UNIVERSITY

Michael Ruth, Ph.D.
mruth@roosevelt.edu
18

Deadlocks

- The cruel irony of programming with threads is that:
 - we want concurrency, so we add threads, but then we get race conditions
 - We don't want race conditions, so we add mutexes, but then we get deadlocks
- **Deadlocks** are a situation in which two threads are waiting for the other to finish and thus neither does
 - Programs don't crash (or even appear to)
 - They just hang...

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
16

Deadlock Avoidance

- Avoiding deadlocks is important and the only consistent, safe means to do so is by designing locking into your threaded application *carefully*
- For example, one particular kind of deadlock is the **deadly embrace**
 - One thread acquires mutex A followed by mutex B and another thread acquires mutex B followed by mutex A
 - Under a timing scenario where they both acquire their first locks before their second locks, neither will ever acquire the second lock
 - Fixing this requires clear rules to assure that this never happens
 - In this case, make it so that all locks are acquired in the same order

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
20

POSIX Threads (pthreads)

- The Linux kernel only provides the underlying primitives that enable threading, the bulk of threading libraries is in user space
 - POSIX standardized a threading library and developers call this standard POSIX Threads or pthreads

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
21

Linux Threading Implementations

- **Linux Threads**
 - Original implementation providing 1:1 threading
 - Designed for a kernel with little support for threading (just clone())
 - Implements pthreads using existing UNIX interfaces
 - Scaled poorly and wasn't in compliance
- **Native POSIX Thread Library (NPTL)**
 - Provides 1:1 threading based around the clone() system call and the kernel's model that threads are just like processes except they share resources
 - Capitalizes on system calls new to newer kernels
 - Improves scalability and conformance issues of Linux Threads

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
22

Pthread API

- The pthread library is very large but contains everything you need to build a multithreaded application
 - Everything is in <pthread.h>
 - Every function begins with pthread_
- All functions can be broken into two main categories:
 - Thread management
 - Synchronization
- If we're going to use pthreads, we need to link it in since it is in a separate library
 - Automated with **-pthread** flag with **gcc**
 - EX: **gcc -o runme -pthread p1.c**

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
23

Creating Threads

- **#include <pthread.h>**
- **int pthread_create (pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine) (void *), void *arg);**
- Upon success, a new thread is created and begins executing the function provided by start_routine passed the sole argument arg
- The function will store the thread ID in the pthread_t if not NULL
- The pthread_attr_t is used to change the default behavior of newly created thread
 - Usually just NULL (we'll skip)
- start_routine must have the following signature:
 - **void * start_thread (void *arg)**

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
24

Thread IDs (TID)

- Thread ID is the thread analogue to PID
 - It does not have to be an arithmetic type
- Retrieving your TID:
 - `#include <pthread.h>`
 - `pthread_t pthread_self(void)`
- Comparing:
 - `#include <pthread.h>`
 - `int pthread_equal(pthread_t a, pthread_t b)`
 - If equal, returns a nonzero value

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
76

Terminating Threads

- Threads may terminate under several circumstances:
 - If a thread returns from start_routine
 - If a thread invokes `pthread_exit()`
 - If a thread is cancelled by another thread via the `pthread_cancel()` function
- All of the threads terminate in the following:
 - The process returns from the main function
 - The process terminates via `exit`
 - The process executes a new binary via `exec`

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
76

Terminating Threads (cont)

- Terminating yourself:
 - `#include <pthread.h>`
 - `void pthread_exit(void *retval)`
- Terminating others:
 - `#include <pthread.h>`
 - `int pthread_cancel(pthread_t thread)`

CSF 35/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
77

Cancellable?

- Whether and when a thread is cancellable depends on its state & type
 - Cancellable state:
 - Either enabled (default) or disabled (puts cancel requests on hold)
 - `#include <pthread.h>`
 - `int pthread_setcancelstate(int state, int *oldstate)`
 - Cancellable type:
 - Either asynchronous or deferred (default)
 - In asynchronous, thread can be killed at any point after cancellation request
 - In deferred, thread can only be killed at cancellation points which represent safe points in pthreads library (critical regions)
 - We can change it:
 - `#include <pthread.h>`
 - `int pthread_setcanceltype(int type, int *oldtype)`

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
26

Joining Threads

- **Joining** allows one thread to wait for another thread to finish (terminate)
 - `#include <pthread.h>`
 - `int pthread_join(pthread_t thread, void **retval)`
 - Upon call, the calling thread waits for the indicated thread
 - Once it terminates, the waiting thread is woken up and if `retval` is not null, provided the return from `pthread_exit`
 - We say that the threads are joined!

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
27

Detaching Threads

- By default, threads are created **joinable**, but we may have them **detach**
 - Which renders them no longer joinable
- Since threads consume resources until joined, threads you do not intend to join should be detached
 - Good practice to explicitly detach or join all threads
- Actual call:
 - `#include <pthread.h>`
 - `int pthread_detach(pthread_t thread)`

CST 357/457 Systems Programming
Threading
Reading: Chapter 7



Michael Ruth, Ph.D.
mruth@roosevelt.edu
28

Pthread mutexes

- For all their power, they are relatively easy to use:
 - Initializing mutexes:
 - `pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;`
 - Locking Mutexes:
 - `#include <pthread.h>`
 - `int pthread_mutex_lock(pthread_mutex_t *mutex)`
 - Unlocking Mutexes:
 - `#include <pthread.h>`
 - `int pthread_mutex_unlock(pthread_mutex_t *mutex)`

Summary

- Discussed thread concepts including terminology, benefits, and costs
- Explained threading models and patterns
- Discussed concurrency and parallelism
- Explained race conditions and synchronization using mutual exclusion
- Discussed deadlock and its avoidance
- Explained pthreads API including thread implementations and using the API

Questions?


