

IPC & Signals

CST 357/457 – Systems Programming

Michael Ruth, Ph.D.

Associate Professor

Computer Science & I.T.

mruth@roosevelt.edu



Objectives

- Discuss signals in terms of purpose, use, lifecycle and their symbolic identifiers
- Explain a common set of signals including their default operations and events
- Discuss basic signal management including sending signals, catching signals, ignoring, and waiting for signals
- Explain reentrancy as it relates to signals and discuss blocking and restoring signals

CST 357/457 Systems Programming
Introduction to Signals
Reading: Chapter 35



Michael Ruth, Ph.D.
mruth@roosevelt.edu

IPC

- Inter-process Communication (IPC)
- There are many reasons one process may wish to communicate with another
 - Synchronization of external resources
 - Producer/consumer relationships
 - Server/client... etc

CST 357/457 Systems Programming
Introduction to Signals
Reading: Chapter 35



Michael Ruth, Ph.D.
mruth@roosevelt.edu

IPC Mechanisms

- There are several methods we will consider:
 - Signals (chapter 10)
 - Pipes (really from I/O but we'll discuss)
 - Queues (message passing)
 - Semaphores
 - Sockets

Signals?

- Signals are software interrupts that provide a mechanism for asynchronous events
 - These events can come:
 - from outside the system
 - EX: User types [CTRL-C] to interrupt processing
 - From activities within a program
 - Divide by zero
- Signals are a primitive form of IPC
 - IPC = InterProcess Communication
- Events occur asynchronously and the program handles them asynchronously
 - Signal **handlers** are **registered** with the kernel

Signal Lifecycle

- First, a signal is **raised**
 - AKA **sent**, **generated**
- The kernel then **stores** the signal
- Finally, when appropriate, the kernel **handles** the signal
 - The kernel can perform one of three functions:
 - Ignore
 - No action is taken
 - Catch & Handle
 - Goes to registered function
 - Perform the Default Action
 - Depends on the signal

Signal Identifiers

- Every signal has a symbolic name that starts with the prefix SIG
 - Defined in header named <signal.h>
 - Actually **int**, but always use symbolic name
- You can generate a list of signals on your system by typing `kill -l`
 - NOTE: `kill` is used to send signals to programs
 - We'll see this later

CSF 357/457 Systems Programming
Introduction to Signals
Reading: Chapter 38



Michael Ruth, Ph.D.
mruth@roosevelt.edu

7

Signals (1)

Signal	Description	Default
SIGALRM	Sent by alarm()	Terminate
SIGCHLD	Child has terminated	Ignored
SIGFPE	Arithmetic Exception	Terminate w/CD
SIGILL	Process Tried to Execute an Illegal Instruction	Terminate w/CD
SIGIO	Asynchronous IO Event	Terminate/Ignore

CSF 357/457 Systems Programming
Introduction to Signals
Reading: Chapter 38



Michael Ruth, Ph.D.
mruth@roosevelt.edu

8

Signals (2)

Signal	Description	Default
SIGPROF	Profiling Timer Expired	Terminate
SIGQUIT	User generated the quit character [CTRL-\\]	Terminate w/CD
SIGSTOP	Suspends Execution of the process	Stop
SIGTERM	Catchable Process Termination	Terminate
SIGTSTP	User Generated Suspend Character [CTRL-Z]	STOP

CSF 357/457 Systems Programming
Introduction to Signals
Reading: Chapter 38



Michael Ruth, Ph.D.
mruth@roosevelt.edu

9

Signals (3)

Signal	Description	Default
SIGURG	Urgent IO Pending	Ignored
SIGUSR2	Process Defined Signal	Terminate
SIGWINCH	Size of controlling terminal window changed	Ignored
SIGXFSZ	File resource limits were exceeded	Terminate w/CD

Basic Signal Management

- `#include <signal.h>`
- `typedef void (*signalhandler_t) (int)`
- `signalhandler_t signal(int signo, sighandler_t handler)`
- Successful call removes the current action for the given signal and handles the signal with the given handler
- Signal Handler:
 - `void my_handler(int signo)`
- You can also use this function to instruct the kernel to IGNORE or go back to the default handler using special values for handler:
 - SIG_DFL – default
 - SIG_IGN – ignore

Waiting For a Signal

- Useful for debugging purposes, the following mechanism blocks until a signal is received that is handled or terminates
 - `#include <signal.h>`
 - `int pause(void)`

Execution & Inheritance

Signal Behavior	Across Forks	Across Execs
Default	Inherited	Inherited
Pending	Not Inherited	Inherited

Mapping Signal Numbers to Strings

- It's often important to get the name rather than the number, so...
 - `extern const char * const sys_siglist[]`
 - Actually your best bet!
 - OR
 - `#include <signal.h>`
 - `void psignal(int signo, const char *msg)`

Sending a Signal

- The `kill()` system call is the basis of the `kill` utility sends a signal from one process to another:
 - `#include <sys/types.h>`
 - `#include <signal.h>`
 - `int kill(pid_t pid, int signo)`
- In normal use, sends `signo` to process identified by `pid`
 - If `pid = 0`, sends `signo` to all processes in process group
 - If `pid = -1`, sends `signo` to all processes it has permission to send a signal to except itself/init
- On success, returns 0, otherwise returns -1 and sets `errno` to one of the following:
 - `EINVAL` (bad signal)
 - `EPERM` (don't have permission)
 - `ESRCH` (bad process or a zombie)

Sending Signals with Kill Utility

- **kill [-signal] pid**
 - Sends TERM by default!
- **kill -9 pid**
 - “Guarantees” that the process will die
- Terminal Signals:
 - <CTRL-C> → SIGINT
 - CTRL-\ → SIGQUIT
 - CTRL-Z → SIGSTP

CSF 35/457 Systems Programming
Introduction to Signals
Reading: Chapter 36



Michael Ruth, Ph.D.
mruth@roosevelt.edu
16

Permissions

- In order to be able to send a signal from one process to another we need permission
 - A process with CAP_KILL (usually root) can send a signal to any process
 - A process without CAP_KILL requires that the effective or real UID equal to the real or saved UID of the receiving process
 - A user can only send signals to processes it owns

CSF 35/457 Systems Programming
Introduction to Signals
Reading: Chapter 36



Michael Ruth, Ph.D.
mruth@roosevelt.edu
17

Sending a Signal to Yourself

- The raise function allows you to send a signal to yourself
 - `#include <signal.h>`
 - `int raise(int signo)`

CSF 35/457 Systems Programming
Introduction to Signals
Reading: Chapter 36



Michael Ruth, Ph.D.
mruth@roosevelt.edu
18

Reentrancy

- A signal can come from anywhere while the process is executing code.
 - It could be in the middle of something delicate that would leave the system in an inappropriate state
- Signal handlers should never assume they know where or what the system was doing and should be very careful about:
 - Manipulating global variables (or use)
- What about system calls that use buffers? or use files? Allocate memory?
 - Some functions are clearly not reentrant
 - If a signal arrives in the middle of a nonreentrant operation, and the signal handler invokes the nonreentrant code, chaos...
- A **reentrant function**, then, is a function that does not manipulate static data, must manipulate only stack data or data supplied by the call, and must NOT call nonreentrant functions

CSF 35/457 Systems Programming
Introduction to Signals
Reading: Chapter 38



Michael Ruth, Ph.D.
mruth@roosevelt.edu
26

Signal Handlers & Reentrancy

- Since we don't know anything about the code that sent us to the handler, we must only ensure that the handler itself does not use functions that are not reentrant
 - There is a list on page 349-350
- We'll learn to block signals to ensure that we don't receive signals at critical times
 - Critical times though depend on reentrancy

CSF 35/457 Systems Programming
Introduction to Signals
Reading: Chapter 38



Michael Ruth, Ph.D.
mruth@roosevelt.edu
27

Signal Sets

- A datatype which represents multiple signals
- `#include <signal.h>`
 - `int sigemptyset(sig_set_t *set);`
 - Sets the set to an empty set (init)
 - `int sigfillset(sig_set_t *set);`
 - Initializes the signal set so that all signals are included
 - `int sigaddset(sig_set_t *set, int signo);`
 - Adds a signal to the set
 - `int sigdelset(sig_set_t *set, int signo);`
 - Deletes a signal from the set
 - `int sigismember(const sig_set_t *set, int signo);`
 - Determines if signal exists in set

CSF 35/457 Systems Programming
Introduction to Signals
Reading: Chapter 38



Michael Ruth, Ph.D.
mruth@roosevelt.edu
28

Signal mask

- The signal mask of a process is the set of signals that are *blocked* from delivery to that process
- `#include <signal.h>`
- `int sigprocmask(int how, const sigset_t *restrict set, sigset_t *restrict oset)`
 - `oset` is a non-null pointer, the current signal mask is returned through it
 - if `set` is non-null, what happens is dependent on `how` variable which could be one of:
 - `SIG_BLOCK` – block the signals on the list
 - `SIG_UNBLOCK` – unblock the signals on the list
 - `SIG_SETMASK` – set the mask to the list

Pending Signals

- To retrieve a list of signals pending for this process:
 - `int sigpending(sigset_t *set)`
 - The result is stored in `set`

Waiting for a Set of Signals

- Allows a process to temporarily change its signal mask and then wait until a signal is raised that either terminates or is handled
 - `#include <signal.h>`
 - `int sigsuspend(const sigset_t *sigmask)`
 - If a signal terminates the process, call never returns
 - If a signal is raised and handled, call returns -1
- A common use of `sigsuspend` is to retrieve signals that might have arrived but were blocked during a critical region
 - Process uses `sigprocmask` to block, then `sigsuspend`

Other Forms of IPC

- Signals are an
 - Event notification mechanism (only notifies other processes of state of this process)
- The other forms of IPC differ in that they:
 - Can actually communicate rather freely (send information back and forth at will)

Pipes

- Pipes are the oldest form of UNIX IPC
 - They have two limitations:
 - half-duplex communication
 - Can only be used between processes that share a common ancestor
 - Pipes are still the most commonly used form of IPC
- There are really two mechanisms for dealing with Pipes in C
 - Formatted Pipes (we only care about this one)
 - Low Level Pipes

Formatted Pipes

- **FILE *popen(char *command, char *mode)**
 - Executes the function specified by command.
 - It creates a pipe between the *calling program and the executed command*, and returns a pointer to a stream that can be used to either read from or write to the pipe.
 - ensures that any streams from previous calls that remain open in the parent process are closed in the new child process

Formatted Pipes and Direction

- If the file mode is "r"
 - The stdout of cmd string is the "input"
 - The file pointer (fp) reads the "output"
- If the file mode is "w"
 - The stdin of cmd string is the "output"
 - The file pointer (fp) writes the "input"

Closing a Pipe

- `int pclose(FILE *stream);`
 - Closes stream opened by `popen`
 - Waits for the command to terminate
 - Returns termination status (or -1 on failure)

popen example 1

```
int main(void) {
    int cnt;
    FILE *pipe_fp;
    char *strings[5] = {"echo", "bravo", "alpha",
                       "charlie", "delta"};

    if ((pipe_fp = popen("sort", "w")) == NULL) {
        perror("popen"); exit(1);
    }
    for(cnt=0; cnt<MAXSTRS; cnt++) {
        fputs(strings[cnt], pipe_fp);
        fputc('\n', pipe_fp);
    }
    pclose(pipe_fp);

    return(0);
}
```

popen example 2

```
int main(void) {
    FILE *pipein_fp, *pipeout_fp;
    char readbuf[80];

    if (( pipein_fp = popen("ls", "r")) == NULL) {
        perror("popen"); exit(1); }
    if (( pipeout_fp = popen("sort", "w")) == NULL) {
        perror("popen"); exit(1); }

    /* Processing loop */
    while (fgets(readbuf, 80, pipein_fp))
        fputs(readbuf, pipeout_fp);

    /* Close the pipes */
    pclose(pipein_fp);
    pclose(pipeout_fp);

    return(0); }
```

Final Note

- Since popen() uses the shell all shell expansion characters are available for use!
- In addition, more advanced techniques such as redirection can be used:
 - `popen("ls ~scottb", "r");`
 - `popen("sort > /tmp/foo", "w");`
 - `popen("sort | uniq | more", "w");`

Summary

- Discussed signals in terms of purpose, use, lifecycle and their symbolic identifiers
- Explained a common set of signals including their default operations and events
- Discussed basic signal management including sending signals, catching signals, ignoring, and waiting for signals
- Explained reentrancy as it relates to signals and discuss blocking and restoring signals

Questions?



CSF 357/457 Systems Programming
Introduction to Signals
Readings: Chapter 20

R ROOSEVELT
UNIVERSITY

Michael Ruth, Ph.D.
mruth@roosevelt.edu
26
