
Introduction to Sockets in C

CST 357/457 – Systems Programming

Michael Ruth, Ph.D.

Associate Professor

Computer Science & I.T.

mruth@roosevelt.edu



Objectives

- Discuss sockets concepts including IPC
- Explain socket creation including domains, types, protocols, and addresses
- Discuss socket functions including send/receive, close, and IO relation
- Explain binding sockets to addresses, listening for connections, making connections, and transferring data
- Discuss client/server operation using sockets

Introduction

- Sockets are the ultimate form of IPC
 - Network IPC
 - Goals was to achieve the same interface for
 - Intermachine communication
 - Intramachine communication
 - Sockets are abstractions of a communication endpoint
 - Socket Descriptors
 - Handle to a socket
 - Very much like file descriptors
 - In fact, the majority of functions that work with files work with sockets

Sockets

- To create a socket we use the `socket` function:
 - `#include <sys/socket.h>`
 - `int socket(int domain, int type, int protocol)`
 - Protocol is usually zero to select the default protocol for the given domain and socket type
 - Socket Type determines the type of communication mechanism
 - Socket Domain determines the nature of the communication

Domains & Types

- Socket Domains:
 - **AF_INET**
 - AF_INET6
 - AF_UNIX
- Socket Type:
 - SOCK_DGRAM
 - UDP interface
 - SOCK_RAW
 - Datagram interface to IP
 - SOCK_SEQPACKET
 - Fixed-length, reliable, sequenced, connection-oriented messages
 - **SOCK_STREAM**
 - TCP interface

Sockets & Functions

- socket and open are very related
- Some important functions and their operations
 - close
 - deallocates the socket
 - read
 - equivalent to recv w/o flags
 - write
 - equivalent to send w/o flags

Sockets and Bidirectional I/O

- Communication via Sockets is bi-directional
- We can disable the I/O on a socket with the following:
 - `int shutdown(int sockfd, int how)`
 - How parameter–
 - SHUT_RD – disables reading
 - SHUT_RW – disables writing
 - SHUT_RDWR – disables both
- Why would we wish to use this:
 - Close deallocates the endpoint only when the last active reference is closed, but shutdown allows us to deactivate the socket
 - Sometimes it is convenient to shut down I/O in a direction to inform the other side that we are done (in that way)

Byte Ordering

- Little/Big Endian
 - Big-endian – the highest byte address occurs in the least significant byte (LSB)
 - EX: Solaris 9 – Sun Sparc
 - Little-Endian – the least significant byte (LSB) contains the lowest byte address
 - EX: Linux 2.4.22 - Intel Pentium
- Host/Network
 - Host implies the computer you're on
 - Network refers to the fact that TCP uses big-endian

Host/Network Byte Order Translation

- `#include <arpa/inet.h>`
- `uint32_t htonl(uint32_t hostint32)`
- `uint16_t htons(uint16_t hostint16)`
- `uint32_t ntohl(uint32_t netint32)`
- `uint16_t ntohs(uint16_t netint16)`

Address Format

- A set of structs allow us to perform 90% of what we need to do
 - **sockaddr**
 - **sockaddr_in**
 - **in_addr**

struct sockaddr

```
struct sockaddr {  
    unsigned short sa_family; // address family, AF_***  
    char sa_data[14]; // 14 bytes of protocol address  
};
```

- This structure holds socket address information for many types of sockets:
 - *sa_family*
 - AF_INET
 - *a_data* contains a ***destination address*** and ***port number*** for the socket

struct sockaddr_in

```
struct sockaddr_in {  
    short int sin_family; // Address family unsigned short  
    int sin_port; // Port number  
    struct in_addr sin_addr; // Internet address  
    unsigned char sin_zero[8]; // Same size as struct  
sockaddr  
};
```

- programmers created a parallel structure for sockaddr – can be substituted
 - **sin_port** – must be in network byte order
 - **sin_addr** – must be in network byte order

IP Address and port in NBO

- A simple and quick mechanism to perform this is:
 - `in_addr_t inet_addr(const char *cp) ;`
 - converts an IP address in numbers-and-dots notation (cp) into an unsigned long
 - Already in Network Byte Order
 - Port must be converted
 - EX: `htons(13) ;`

ntoa (network to address)

- `char *inet_ntoa(struct in_addr in);`
 - Converts the address specified by `in_addr` to a human readable string in numbers-and-dots notation

Address Resolution

- Ideally, we shouldn't be forced to know the internal structures
 - Additionally, we wish to be able to perform hostname and server lookups
 - w/o scanning /etc/hosts or /etc/services
- **struct hostent *gethostbyname(char *hostname) ;**
 - Resolves the IP name for the given hostname

struct hostent

```
struct hostent {  
    char* h_name;  
    char** h_aliases;  
    int h_addrtype;  
    int h_length;  
    char** h_addr_list;  
    #define h_addr h_addr_list[0]  
};
```

- h_name: This is the official name of the host, i.e. the full address.
- h_aliases: a pointer to the list of aliases (other names) the host might have.
- h_addrtype: The type of address this host uses.
- h_length: The length of the address. Different address types might have different lengths.
- h_addr_list: A pointer to the list of addresses of the host. Note that a host might have more than one address, as explained earlier.
- h_addr: In older systems, there was only the h_addr field, so it is defined here so old programs could compile without change on newer systems.

Associating Addresses With Sockets

- `int bind(int sockfd, struct sockaddr *addr, socklen_t len)`
 - Returns 0 if ok, -1 on error
 - Some restrictions:
 - Address specified must be valid for the machine this is running on
 - Address must match the format specified by address family
 - The port number cannot be less than 1024 w/o appropriate permissions
 - Only one socket per address (there are exceptions)
 - **portnumber == 0**
 - Chooses the first unused port
 - **INADDR_ANY**
 - automatically fill in the IP address of the machine the process is running on

Defaults and Discovery

- Since we have some default modes, let's discuss how we can find out what we actually are bound to
 - `int getsockname(int sockfd, struct sockaddr *addr, socklen_t alenp)`
 - Discover the address bound to the socket
 - `int getpeername(int sockfd, struct sockaddr *addr, socklen_t len)`
 - If the socket is connected, return the address of its peer

Time to actually Connect

- If we are dealing with connection-oriented network service (SOCK_STREAM or SOCK_SEQPACKET) we need to establish a connection
 - `int connect(int sockfd, struct sockaddr *addr, socklen_t len)`
 - There may be many reasons for failure
 - Additionally, can connect be used with SOCK_DGRAM ?
 - Yes, but....

Server Mechanisms

- A server announces that it is willing to accept connect requests by calling:
 - `int listen(int sockfd, int backlog)`
 - Backlog is a queue size
- We retrieve the connections from the server by calling
 - `int accept(int sockfd, struct sockaddr *addr, socklen_t *len)`
 - The descriptor returned is the descriptor connected to the process which called connect
 - Accept blocks!

Transferring data

- Connection Oriented
 - Sock presented by sockfd must have already been opened and connected
 - send
 - recv
- Connectionless
 - No need to open connection, simply use the following methods:
 - sendto
 - recvfrom

send

- `ssize_t send(int sockfd, void *buf, size_t nbytes, int flags)`
- Identical to write except for flags argument:
 - MSG_DONTROUTE
 - MSG_EOR
 - MSG_DONTWAIT

recv

- `ssize_t recv(int sockfd,
void *buf, size_t nbytes,
int flags)`

– Again same as read except for flags:

- `MSG_PEEK`
- `MSG_WAITALL`
- `MSG_TRUNC`

sendto & recvfrom

- `ssize_t sendto(int sockfd, void *buf, size_t nbytes, int flags, struct sockaddr *destaddr, socklen_t destlen)`
- `ssize_t recvfrom(int sockfd, void *buf, size_t nbytes, int flags, struct sockaddr *addr, socklen_t addrlen)`

Socket Options

- We can get and set socket options
 - Three kinds of options:
 - Generic options for all socket types
 - Options that are managed at the socket level, but depend on the underlying protocols for support
 - Protocol-specific options unique to each individual protocol
 - `int setsockopt(int sockfd, int level, int option, void *val, socklen_t len)`
 - `int getsockopt(int sockfd, int level, int option, void *val, socklen_t lenp)`
 - Level identifies the protocol
 - If generic then level == SOL_SOCKET

Socket Options Table

Options	Description
SO_ACCEPTCONN	Return whether a socket listening
SO_BROADCAST	Broadcast datagrams if val is nonzero
SO_DONTROUTE	Bypass normal routing
SO_KEEPALIVE	Periodic keep alive messages enabled if val is nonzero
SO_REUSEADDR	Reuse addresses in bind if val is nonzero
SO_SNDBUF	The size in bytes of send buffer
SO_RCVBUF	The size in bytes of receive buffer

Client/Server Model

- The client-server model is used to divide the work of Internet programs into two parts:
 - Server:
 - Passive (slave)
 - Waits for requests
 - Upon receipt of requests, processes them and then serves replies
 - Client:
 - Active (master)
 - Sends requests
 - Waits for and receives server replies

Client Psuedocode

```
get the server's address
form a working address that can be used to talk over
Internet.

connect to the server

while (!done) do:
    wait until there's either information from the server,
    or
        from the user.

    If (information from server) do
        parse information
        show to user, update local state information, etc.

    else {we've got a user command}
        parse command
        send to server, or deal with locally.

done
```

Server pseudocode

```
bind a port on the computer, so Clients will be able
to    connect
forever do:
    listen on the port for connection requests.
    accept an incoming connection request
    if (this is an authorized Client)
        while (connection still alive) do:
            receive request from client
            handle request
            send results of request, or error messages
        done
    else
        abort the connection
done
```

Summary

- Discussed sockets concepts including IPC
- Explained socket creation including domains, types, protocols, and addresses
- Discussed socket functions including send/receive, close, and IO relation
- Explained binding sockets to addresses, listening for connections, making connections, and transferring data
- Discussed client/server operation using sockets

Questions?

